

CONSILIUM EXECUTION BLUEPRINT

v9.1 (MVP)

Target: Bittensor SN4 (Optimization / Dense QUBO)

MVP Deadline: 28 Feb 2026

Status: Engineering Specification (for implementation by 2-person team + AI)

Confidentiality: Internal / do not distribute

0. Executive Summary

This document is an implementation-focused translation of the “E8 + Topology + Quantum Scout” approach for Bittensor SN4.

Core idea (what to build): - SN4 provides **Dense QUBO** instances Q under a strict time window. - We extract fast **topological signals** from the task’s interaction structure (“holes/loops” → invariants like β_1) using an IBM/Qiskit **Quantum Scout** with strict timeouts and hard fallbacks. - We construct a deterministic **8D topological latent space** and a projection matrix $W \in R^{(8 \times N)}$ (“topological principal components”). - We **quantize** the latent space against the **E8 root system** (240 roots = vertices of the Gosset polytope 4_21) by snapping to the nearest E8 root vector. - We **lift** the chosen E8 root direction back to an N-bit seed via W^+ , then run a short **Lift & Repair** local search on the original Q to produce a valid submission.

MVP goal: - Implement the full SN4 miner loop on **testnet** and demonstrate non-random positive incentive/score under strict latency constraints.

What is intentionally *not* explained here: - Philosophical / “why E8” motivation. The spec focuses on I/O and determinism so an AI can implement it.

1. MVP Definition (Scope)

1.1 Goal

Achieve stable positive incentive/score on SN4 **testnet** by running a single miner instance continuously and submitting valid solutions.

1.2 Success Metrics (KPIs)

Define these numbers before Week 2 and keep them stable for the month:

- **Median incentive/score $\geq S_{\text{target}}$** over a continuous run of T_{eval} hours.
- **p95 end-to-end latency $\leq T_{\text{p95}}$** seconds (job received → tx submitted).
- **Acceptance rate $\geq R_{\text{accept}}$** (valid solutions / total attempts).
- **Stability:** $\geq 24\text{h}$ continuous operation without crash.

Recommended initial targets (edit after first baseline run):

- $S_{\text{target}} = 0.001$
- $T_{\text{p95}} = 45\text{s}$
- $R_{\text{accept}} = 80\%$
- $T_{\text{eval}} = 24\text{h}$

1.3 Non-Goals (MVP)

- No GUI/front-end (CLI only).
- No autoscaling / multi-node cluster management.
- No mainnet risk in February (testnet only).
- No “enterprise” observability; keep only minimal structured logs.

1.4 Strategic Context (Sniper Doctrine)

SN4 is treated as a **time-windowed optimization market**. The intended niche is:

- Dense / high-connectivity QUBOs where naive GPU heuristics struggle with state update costs.
- We ignore low-complexity instances and focus on tasks where the E8+topology pipeline gives asymmetric advantage.

2. The Interface: Bittensor SN4 Contract (Must-Be-True)

Important: v9.0 contained placeholder method names. For v9.1, treat SN4 I/O as **contract-first**: verify the real Synapse schema and submission call in the actual SN4 codebase before coding.

2.1 Task Source

Miner receives tasks from SN4 validators via a Bittensor Synapse.

MVP deliverable: - A single function that converts “Raw Synapse” → (Q, meta) where: - Q: dense matrix ($N \times N$, float32 preferred) - meta: includes at minimum job_id, deadline_ts or timeout_s, and any validator-provided scoring hints.

2.2 Task Data Model (Schema to confirm)

Minimum fields we assume exist (or can be derived): - N (int): number of variables (expected 2000–5000). - Q (matrix): QUBO weights, typically in [-1, 1] (confirm). - domain (enum): variables are either {0,1} or {-1,1} (confirm; implement conversion). - time_window_s (float): hard deadline for response (confirm).

2.3 Solution Output Model

Miner must output: - solution: list/array of length N (binary/spins per subnet contract) - optional energy (float) if accepted by protocol (confirm) - signature / wallet metadata as required by Bittensor

2.4 Scoring Objective (MVP assumption)

Assume validators reward: - Lower energy $H(s)$ is better - Latency is a multiplier/penalty - Invalid format or deadline miss → zero or negative

Define energy in a single canonical form internally, and convert at boundaries: - Spin form: $s \in \{-1, +1\}^N$, $H(s) = s^T Q s$ - Binary form: $x \in \{0, 1\}^N$, $H(x) = x^T Q x$

3. System Architecture (Single-Node MVP)

3.1 Components

All components run on one host (laptop or server), with external calls to IBM/Qiskit and an optional solver service.

- **Ingest:** Synapse → dense Q, meta
- **Quantum Scout:** (Q, meta) → topological features F (with strict timeout + fallback)
- **W Builder:** (Q, F) → $W \in R^{(8 \times N)}$ + cache key
- **E8 Quantizer:** W, Q → candidate roots r_i and/or best r^*
- **Lift & Repair:** (W, r^* , Q) → valid N-bit solution s
- **Submit:** s → on-chain/testnet submission
- **Logger:** writes one JSON line per attempt (minimal observability)

3.2 Degraded Modes (Required for MVP)

The pipeline must always produce *some* output before deadline: -

Degraded-Q: Quantum Scout timed out → use cached F_cache -

Degraded-W: W build failed → use cached W_cache or

deterministic default W_default - **Degraded-S:** Repair cannot finish → return best-so-far seed solution

3.3 Hardware Roles (Conceptual)

- **Orchestrator (Local host):** parsing, caching, building w, lift & repair, submission.
- **Scout (IBM Quantum / Qiskit Runtime):** topological feature estimation under strict timeout.
- **Hammer (optional external solver):** if configured, solves/warms-starts candidate solutions within a time slice.

4. The “Sniper Loop” Pipeline (Deterministic, Deadline-Aware)

4.1 Hard Time Budget (initial)

This is a starting point; tune after first testnet run.

- Step 1 Ingest: 0.5s
- Step 2 Quantum Scout: 3.0s (hard timeout)

- Step 3 Build W: 5.0s
- Step 4 E8 Quantize: 0.2s
- Step 5 Lift & Repair: \leq (deadline - elapsed - submit_budget)
(target 2-20s)
- Step 6 Submit: 1.0s

4.2 Filtering (Sniper Scope)

We intentionally ignore tasks that do not match the dense/high-connectivity niche.

Define **effective density**:

$$\text{density_eps}(Q, \text{eps}) = \text{count}(|Q_{ij}| > \text{eps}) / (N^2)$$

Initial policy: - If $N < 2000$: DROP (optional; tune) - If $\text{density_eps}(Q, \text{eps}=1e-3) < 0.3$: DROP (tune)

4.3 End-to-End Steps

- 1) **Ingest** - Parse synapse, validate shapes/types, convert to internal canonical domain (spin or binary). - Compute cheap stats: N, density_eps, min/max weight, symmetric check.
- 2) **Quantum Scout (timeout hard)** - Compute topological features F: - must include at least beta1_estimate (even if approximate) - may include additional scalars: components_estimate, spectral_near_zero, etc. - If timeout/error: load F_cache keyed by (subnet_version, N_bucket, density_bucket) or last-good.
- 3) **Build W** - Build $W \in R^{(8 \times N)}$ deterministically from (Q, F) (Section 6). - Cache W for reuse across tasks (keyed by N_bucket + beta1_bucket). - Precompute W^+ (pseudo-inverse) when W changes.
- 4) **E8 Quantize** - Generate E8 roots (240 vectors in R^8) deterministically (Appendix A). - Compute reduced matrix $Q8 = W Q W^\top (8 \times 8)$. - Select root candidates r by minimizing reduced energy $E8(r) = r^\top Q8 r$. - Keep top B candidates (beam width; MVP)

start B=8).

- 5) **Lift & Repair (anytime)** For each candidate r in order: - Lift: $s_{\text{seed}} = W^+ r$ (vector in R^N) - Binarize/spinize: $s_0 = \text{binarize}(s_{\text{seed}})$ (Section 6) - Repair: run greedy local search on original Q under a strict time slice. - Track best (energy, solution) and stop when remaining time is low.
 - 6) **Submit** - Submit best solution found. - No artificial delays: submit as soon as ready.
-

5. Mathematical & Data Definitions

5.1 QUBO Canonicalization

Pick one internal representation for MVP (recommended: **spin** {-1,+1}): - If subnet uses {0,1}, convert to spin or implement both energy functions and repair deltas.

5.2 E8 Root System (Gosset Polytope 4_21)

We use the E8 root system as a fixed codebook of 240 vectors in R^8 : - All roots have equal norm ($\|r\|^2 = 2$) - Two families: - Type A: permutations of $(\pm 1, \pm 1, 0, 0, 0, 0, 0, 0) \rightarrow 112$ roots - Type B: $(\pm 1/2, \dots, \pm 1/2)$ with an even number of minus signs $\rightarrow 128$ roots

Engineering note: - In MVP we do **not** need to implement Weyl-group actions explicitly; we only need the deterministic root codebook and snapping.

5.3 Projection Matrix W

Corrected definition (v9.1): - $W \in R^{(8 \times N)}$ maps N-dimensional states into an 8D latent space. - $v = W s$, where $v \in R^8$, $s \in R^N$ (a candidate state/seed). - W is derived from topological structure (Section 6), not hand-tuned.

5.4 Snapping / Quantization

Given $v \in R^8$, snap to nearest E8 root:

$$r^* = \operatorname{argmin}_{\{r \text{ in Roots_E8}\}} \|v - r\|_2$$

Since all roots share a constant norm, this is equivalent to maximizing dot product:

$$r^* = \operatorname{argmax}_{\{r \text{ in Roots_E8}\}} \langle v, r \rangle$$

5.5 Pseudo-Inverse

Given W has full row rank, compute:

$$W^+ = W^T (W W^T)^{-1} \quad \# \text{ dimensions: } (N \times 8)$$

Compute and cache W^+ whenever W changes.

6. E8 Regularizer (Engineering Specification)

6.1 Inputs / Outputs

Inputs: - Q : dense $N \times N$ float32 - F : topological feature vector (at minimum β_1 _estimate)

Outputs: - W : $8 \times N$ float32 - W_{plus} : $N \times 8$ float32 - Roots_E8 : list of 240 vectors in R^8 (cached global constant)

6.2 Building W (MVP-Implementable)

We need 8 stable “topological principal components” for the task graph.

MVP method (deterministic, fast, reproducible): 1) Build a **skeleton graph** G_{skel} from Q : - For each node i , keep the k strongest interactions by $|Q_{ij}|$. - Symmetrize the edge set. - Set $k = \operatorname{clamp}(k_{\text{min}}, k_{\text{max}}, f(\beta_1 \text{ estimate}, N))$. 2) Build sparse

adjacency A and graph Laplacian $L_0 = D - A$. 3) Compute 8 eigenvectors associated with the smallest non-trivial eigenvalues of L_0 : - Stack them as rows of W (after normalization / orthonormalization).

Recommended MVP defaults: - k_min = 32, k_max = 256 - f(beta1, N) = round(16 + 2*min(beta1, 64)) (tune)

Notes: - This matches the “topological components” intent and keeps runtime feasible at $N \approx 5000$. - The Quantum Scout controls beta1_estimate, which controls skeletonization and stability of W.

6.3 Reduced Energy Selection in E8 Space

Compute reduced matrix:

$$Q8 = W Q W^T \quad \# 8 \times 8$$

For each root r:

$$E8(r) = r^T Q8 r$$

Select beam B roots with the lowest $E8(r)$.

6.4 Lift & Repair

For each candidate root r: 1) Lift:

$$s_seed = W^+ r \quad \# N\text{-dimensional real vector}$$

2) Binarize/spinize: - If internal domain is spins: $s0_i = +1$ if $s_seed_i \geq 0$ else -1 - If internal domain is binary: $x0_i = 1$ if s_seed_i

≥ 0 else 0 3) Repair (greedy local search on original Q): - Run until time slice exhausted or no improving flips. - Always maintain best-so-far solution.

Repair must be **anytime**: - If time remains: continue improving - If close to deadline: stop and return current best

Implementation detail for speed: - Maintain $g = Q s$ (or $Q x$) and update g after each flip so each flip is $O(N)$, not $O(N^2)$. (See Appendix B.)

6.5 Why this can work (engineering framing)

The E8 root codebook provides a compact set of highly symmetric directions in 8D. The hypothesis is: - The 8D embedding concentrates “energy landscape shape” into a small space. - Snapping to E8 roots yields seeds that are closer to good minima than random initialization.

This is an empirical claim for MVP: - We validate only by testnet score and golden-set energy distributions.

7. Quantum Scout (Engineering Specification)

7.1 Purpose

Return a **topological feature vector F fast**, under strict timeouts, robust to IBM API issues.

Minimum required output for MVP: - beta1_estimate (float or int) - confidence (0..1) or source (quantum | cache | fallback)

7.2 I/O Contract

Input: - Q (dense $N \times N$) Output: - $F = \{ \text{beta1_estimate}, \dots \}$

7.3 Algorithm Sketch (LGZ-like Rank/Laplacian Estimation)

We estimate β_1 via the kernel dimension of the combinatorial 1-Laplacian $L1$: - Construct a simplicial representation from Q (MVP: use skeleton graph and optional 2-simplices by threshold). - Build boundary operators and $L1$. - Encode $L1$ as a Hermitian operator / Hamiltonian. - Use Qiskit Runtime primitives to estimate spectral

weight near zero.

Heat-trace estimator (conceptual):

$$\beta_1 \approx \text{Tr}(\exp(-\tau L_1)) \quad \text{for large } \tau$$

MVP reality requirement: - The runtime program must return within $T_{\text{quantum}}=3\text{s}$ or we fall back. - If quantum execution cannot meet this, run the same estimator on a smaller sampled subcomplex (document the sampling).

MVP fallback (must be implemented regardless): - If quantum is unavailable, estimate a graph-level cycle count on the skeleton: - For a graph (1-complex), $\beta_1_{\text{graph}} = m - n + c$ (edges - nodes + connected components). - Use this only as a degraded signal for choosing k and stabilizing W .

7.4 Timeouts, Budgets, Fail-Safes

- Hard timeout: $T_{\text{quantum}} = 3.0\text{s}$ (configurable).
- Budget: $\leq Q_{\text{calls_per_hour}}$ (initial 10/h).
- Fail-safe: on timeout or non-2xx response:
- Use cached F_{cache}
- Mark attempt as $\text{degraded}_{\text{quantum}}=\text{true}$

7.5 Cache Keys

Cache F by coarse buckets: - $N_{\text{bucket}} = \text{round}(N / 250) * 250$ - $\text{density}_{\text{bucket}} = \text{round}(\text{density}_{\text{eps}} / 0.05) * 0.05$ - $\text{weight}_{\text{stats}}_{\text{bucket}}$ (optional)

8. External Solver (Optional Accelerator)

MVP allows one of these: - **A)** No external solver: use only Lift & Repair (Section 6.4). - **B)** External solver service (DA / cloud GPU) that accepts a QUBO and returns a candidate solution.

If using an external solver, define a stable interface:

8.1 API Contract (example)

Request:

```
{  
    "job_id": "string",  
    "domain": "spin|binary",  
    "Q": "packed matrix or URL",  
    "time_limit_s": 20.0,  
    "warm_start": "optional seed solution"  
}
```

Response:

```
{
```

```
"solution": [0,1,0,...],  
  
        "energy": -123.45,  
  
        "solver_time_s": 18.7,  
  
        "status": "ok|timeout|error"  
    }  

```

8.2 Failure Policy

- If external solver fails/slow: immediately continue with local Lift & Repair using s_seed.
-

9. Testing & Validation (MVP, Minimal)

9.1 Golden Dataset (local)

Save ~20 real tasks from testnet to tests/data/ as JSON (synapse snapshots).

Required local command: - run_pipeline --input tests/data/golden_01.json --deadline 45s

Acceptance (local): - Produces a correctly shaped solution. - Computes energy without NaN/overflow. - Runs under the deadline budget.

9.2 Testnet Run (the real proof)

Run miner on SN4 testnet for $\geq 2h$ (then 24h). Log and track: - latency per stage - final energy - accepted vs rejected

submissions - score/incentive if available

10. Implementation Roadmap (Feb 2026)

Week 1 — Skeleton

- Implement Ingest + Submit to testnet.
- Implement logging.
- Dummy solver produces valid-shaped random solutions (for protocol verification).

Week 2 — Brain (Quantum + E8 core)

- Implement E8 roots generator + snapping.
- Implement W builder (skeleton graph + spectral embedding).
- Implement Quantum Scout call + strict timeout + cache.

Week 3 — Lift & Repair

- Implement W^+ + lifting + binarize/spinize.
- Implement greedy repair with $O(N)$ flip deltas.
- Golden dataset pipeline.

Week 4 — Tuning & Proving

- Tune time budgets, beam width B, skeleton k, thresholds.
 - 24h testnet run and KPI evaluation.
-

11. Risks & Mitigations (MVP)

- 1) **IBM latency / quota** → strict timeout + cached features + degraded mode.
 - 2) **SN4 protocol changes** → isolate synapse parsing and submission behind interfaces.
 - 3) **Dense matrix performance** → enforce float32 + BLAS + memory budget checks.
 - 4) **Projection error (W bad)** → beam width > 1 + repair randomized restarts.
 - 5) **Correlation gap (E8 reduced energy ≠ real energy)** → evaluate top-B roots + pick best after repair.
-

12. Reference Implementation Outline (Minimal, AI-Friendly)

Recommended Python module boundaries (names are suggestions; keep interfaces stable):

- miner/ingest.py: synapse parsing → (Q, meta)
- miner/quantum_scout.py: Q → F with timeout + cache + fallback
- miner/w_builder.py: (Q, F) → W, W_plus
- miner/e8.py: root generator + snapping + reduced energy ranking
- miner/lift_repair.py: lift, binarize/spinize, greedy repair, energy
- miner/submit.py: submit to subnet
- miner/main.py: sniper loop + budgets + logging

Minimal CLI commands for MVP:

```
- miner run --subnet 4 --network testnet --wallet <name>
- miner replay --input tests/data/golden_01.json --deadline 45s
```

13. Configuration (MVP)

13.1 Environment Variables

- IBM_TOKEN (if using Qiskit Runtime)
- SOLVER_API_KEY / SOLVER_URL (if using an external solver)
- BT_WALLET_NAME / BT_WALLET_HOTKEY (per Bittensor conventions; confirm)

13.2 Runtime Parameters (config file or CLI)

- Timeouts: T_ingest, T_quantum, T_W, T_repair, T_submit
 - Beam width: B
 - Skeletonization: k_min, k_max, eps_density
 - Determinism: RNG_SEED (for repair restarts; default fixed)
-

Appendix A — Deterministic Generator for 240 E8 Roots

Type A (112 roots): - Choose 2 positions out of 8 for non-zero entries. - Assign each of the two entries a sign ± 1 . - All permutations of positions and sign choices.

Type B (128 roots): - All 8-tuples of $\pm 1/2$ with an even number of negative signs.

Implementation note: - Return a stable ordering (lexicographic) for reproducibility.

Appendix B — Greedy Repair: Fast ΔE Updates (Spin Form)

For spins $s \in \{-1, +1\}^N$, energy:

$$H = s^T Q s$$

Maintain $g = Q s$ (vector length N).

Flipping spin i ($s_i := -s_i$) changes energy by:

$$\Delta H_i = -4 * s_i * g_i + 4 * Q_{ii}$$

(Derive/confirm based on the exact diagonal convention used by SN4 tasks; adjust once the dataset is verified.)

After flipping i , update:

$$g := g + (-2 * s_i_{old}) * Q[:, i]$$

This makes each flip $O(N)$.

Appendix C — Minimal Structured Log Schema (JSONL)

One line per task attempt:

```
{  
  
    "ts": "iso8601",  
  
    "job_id": "string",  
  
    "N": 5000,  
  
    "density_eps": 0.92,  
  
    "degraded_quantum": false,  
  
    "beta1": 37.0,  
  
    "beam_B": 8,  
  
    "energy": -123.45,  
  
    "latency_total_s": 31.2,
```

```
"latency": { "ingest": 0.2, "quantum": 2.9, "W": 3.8, "repair": 23.1,  
"submit": 1.0 },  
  
"submit_status": "ok|rejected|timeout|error"  
  
}
```

Appendix D — MUST-CHECK Items (Do Before Coding)

- Confirm SN4 synapse schema and actual submission API in the current subnet repository.
- Confirm variable domain and diagonal conventions in energy computation.
- Confirm whether the validator expects energy or only a solution vector.
- Confirm IBM/Qiskit runtime latency and select an executable strategy that fits T_quantum.